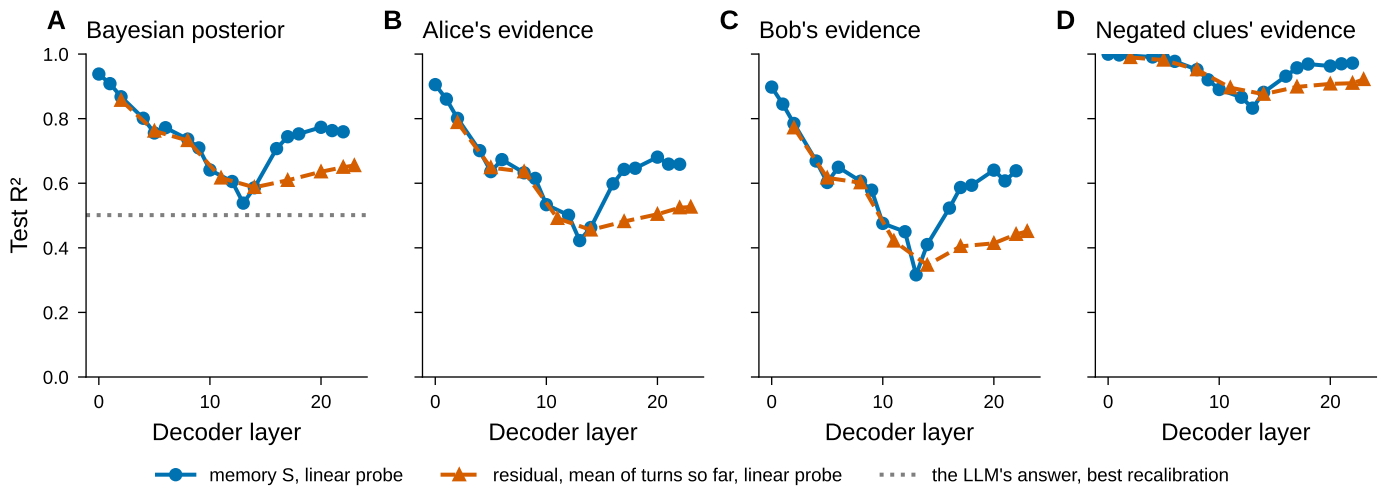
A
Bayesian posterior
B
Alice's evidence
C
Bob's evidence
D
Negated clues' evidence
Test R²
1.0
0.8
0.6
0.4
0.2
0.0
0
10
20
Decoder layer
memory S, linear probe
residual, mean of turns so far, linear probe
the LLM's answer, best recalibration