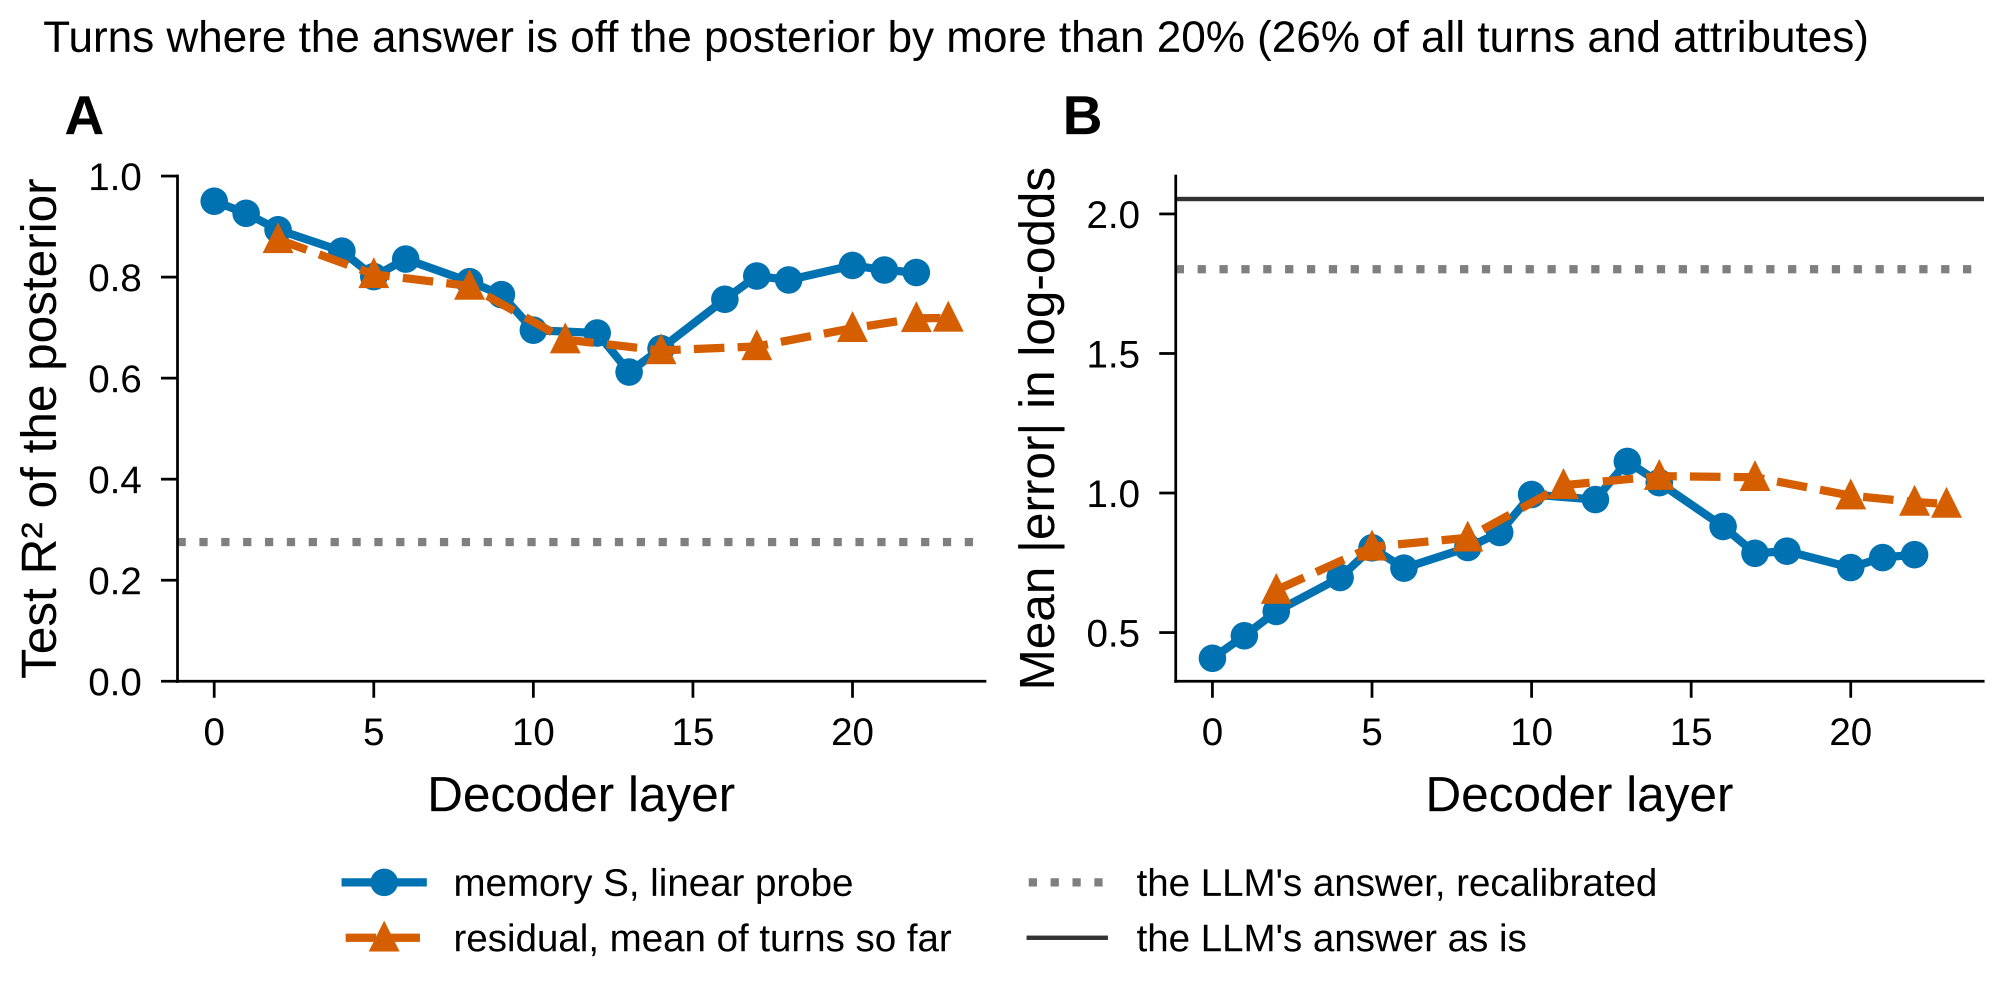
Turns where the answer is off the posterior by more than 20% (26% of all turns and attributes)
A
B
Test R² of the posterior
Mean |error| in log-odds
Decoder layer
Decoder layer
memory S, linear probe
residual, mean of turns so far
the LLM's answer, recalibrated
the LLM's answer as is