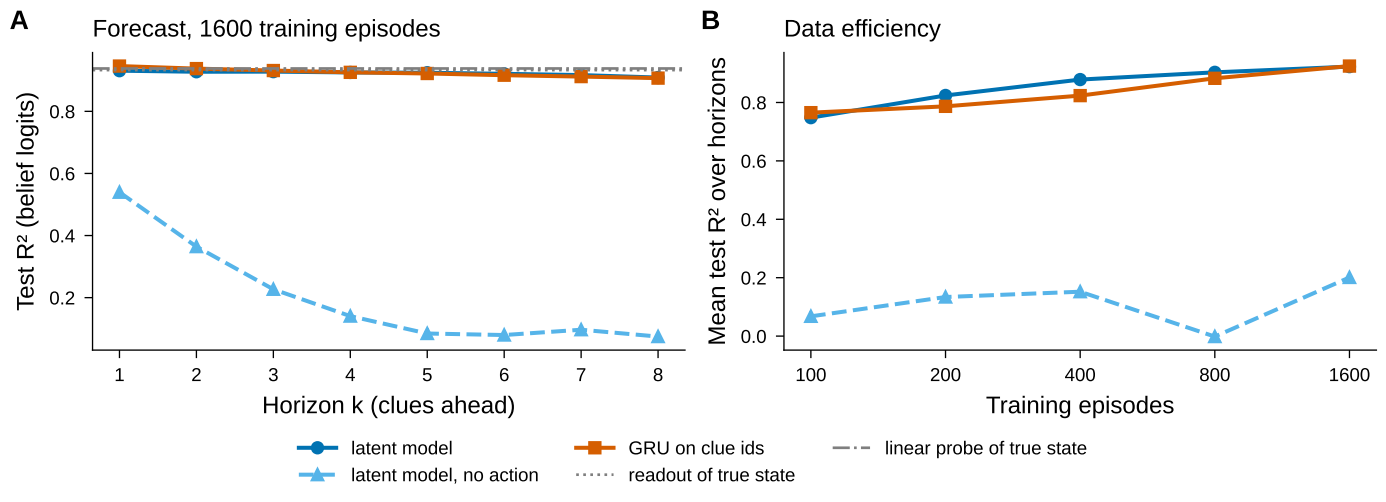
A
Forecast, 1600 training episodes
Test R² (belief logits)
0.8
0.6
0.4
0.2
1
2
3
4
5
6
7
8
Horizon k (clues ahead)
B
Data efficiency
Mean test R² over horizons
0.8
0.6
0.4
0.2
0.0
100
200
400
800
1600
Training episodes
latent model
latent model, no action
GRU on clue ids
readout of true state
linear probe of true state