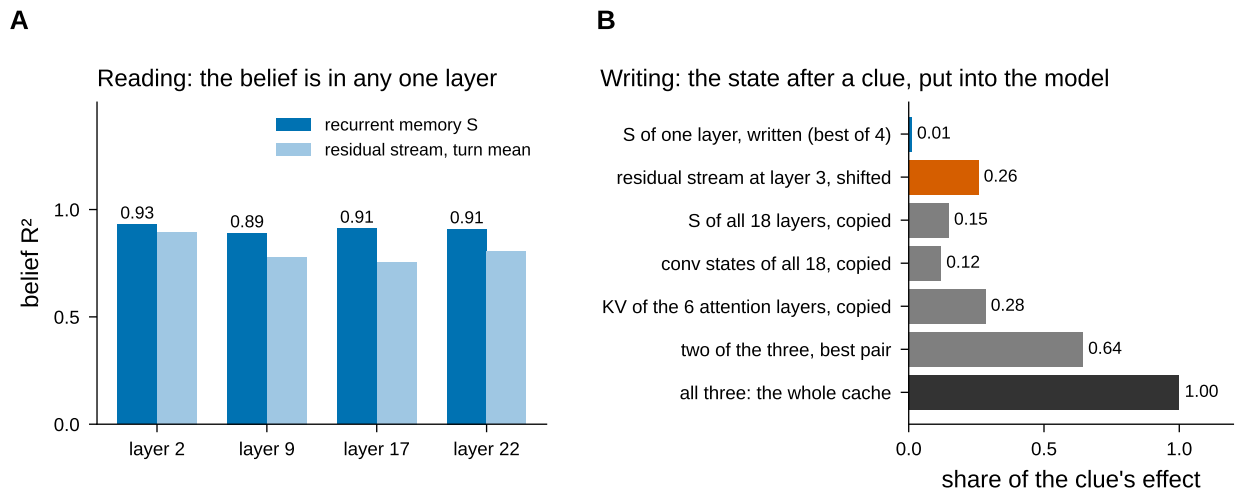
A
Reading: the belief is in any one layer
recurrent memory S
residual stream, turn mean
belief R²
1.0
0.5
0.0
0.93
0.89
0.91
0.91
layer 2
layer 9
layer 17
layer 22
B
Writing: the state after a clue, put into the model
S of one layer, written (best of 4)
0.01
residual stream at layer 3, shifted
0.26
S of all 18 layers, copied
0.15
conv states of all 18, copied
0.12
KV of the 6 attention layers, copied
0.28
two of the three, best pair
0.64
all three: the whole cache
1.00
0.0
0.5
1.0
share of the clue's effect