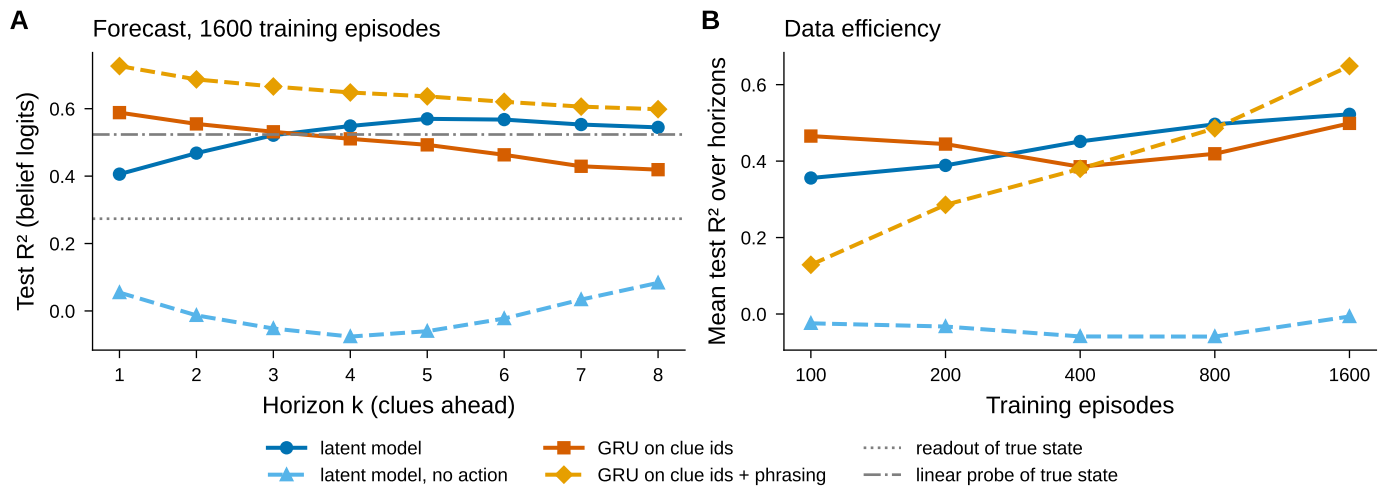
A
Forecast, 1600 training episodes
Test R² (belief logits)
0.6
0.4
0.2
0.0
1
2
3
4
5
6
7
8
Horizon k (clues ahead)
B
Data efficiency
Mean test R² over horizons
0.6
0.4
0.2
0.0
100
200
400
800
1600
Training episodes
latent model
latent model, no action
GRU on clue ids
GRU on clue ids + phrasing
readout of true state
linear probe of true state