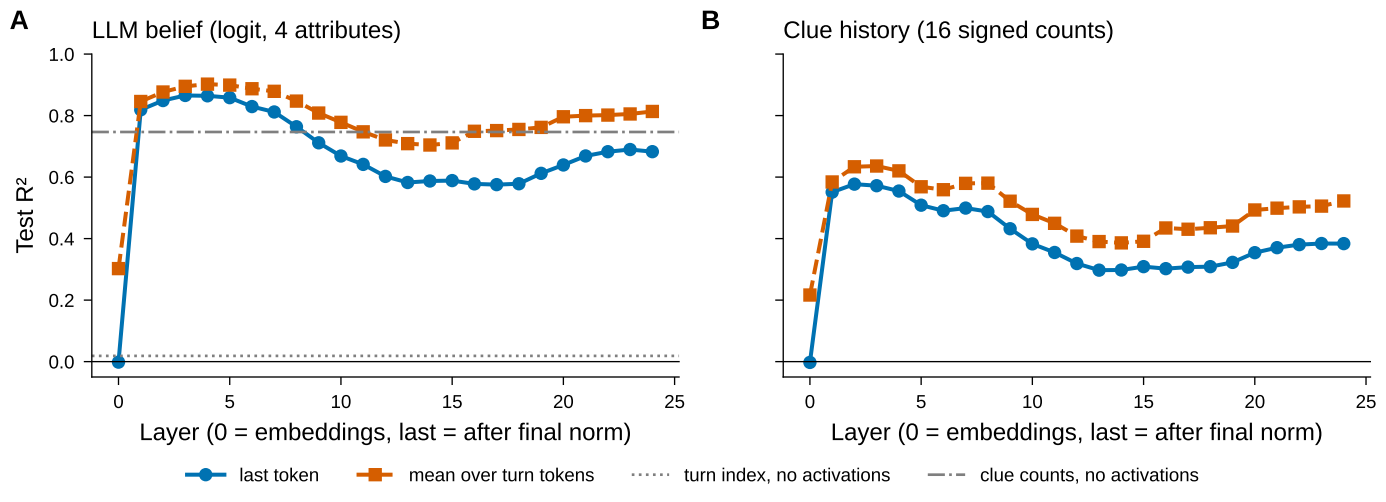
A
LLM belief (logit, 4 attributes)
Test R²
1.0
0.8
0.6
0.4
0.2
0.0
0
5
10
15
20
25
Layer (0 = embeddings, last = after final norm)
B
Clue history (16 signed counts)
0
5
10
15
20
25
Layer (0 = embeddings, last = after final norm)
last token
mean over turn tokens
turn index, no activations
clue counts, no activations